

# Problem 1: Min and Max

Given an array `a[0..n]` of `n` positive integers build a data structure to answer two different types of queries:

1. `update(l, r, t)` that updates every `a[i]` with `min(a[i], t)`, where `l ≤ i ≤ r`
2. `max(l, r)` that returns the largest value in `a[l..r]`

A sequence of `m` queries is provided and the target solution must run in  $O((n + m) \log(n))$  time.

To solve the problem it is provided an implementation of a Segment Tree.

## Segment Tree Implementation

Each `Node n` of the segment tree contains:

- `key: u32`, the maximum of the segment covered by `n`
- `left_edge: usize`, the left edge of the segment covered by `n`
- `right_edge: usize`, the right edge of the segment covered by `n`
- `left_child: Option<Box<Node>>`, the left subtree rooted in `n`
- `right_child: Option<Box<Node>>`, the right subtree rooted in `n`

The tree is represented as a single `Node root`.

## Construction of the Tree

The segment tree is built recursively, passing:

- the array `a` received in input
- the current node `curr_node`
- `0` and `a.len()-1` as the first segment edges

If `l == r` it means that we reached a leaf, and hence we set `curr_node.key = a[l]`. Otherwise, the following operations are done:

- compute `m`, the half of the current node segment
- build a new node `left_child` that covers the segment `[l, m]`
- recursively build the left subtree, using `left_child` as `curr_node`
- build a new node `right_child` that covers the segment `[m+1, r]`
- recursively build the right subtree, using `right_child` as `curr_node`
- compare the `right_child.key` and `left_child.key` and assign the maximum as key to `curr_node`
- add `right_child` and `left_child` as children of `curr_node`

# Max

The first kind of query is of the form `max(l, r)` that returns the largest value in `a[l..r]`.

The function

```
fn max_rec(curr_node: &Node, l: usize, r: usize) -> u32
```

has as parameters:

- `curr_node`, the current root of the tree
- `l`, the left endpoint of the queried segment
- `r`, the right endpoint of the queried segment

and it recursively computes the maximum in `[l, r]`.

The function behaves as follows:

- checks that `l <= r`
- if `curr_node.left_edge == curr_node.right_edge` means that the current node is a leaf, hence returning the current node's key
- if the `curr_node` segment covers exactly the queried one `[l, r]`, hence return the current node's key
- if the queried segment is fully contained on either the `left_child` or `right_child` return the result recursively computed on it
- if the queried segment spans on both the left and right subtree of the current node we compute the middle element `m` and compute the maximum from the left and right by recurring on `[l, m]` and `[m+1, r]` and return the greatest between the two.

**Time Complexity:**  $O(\log(n))$

- the function traverses the segment tree from the root to the leaves, and for each node, it performs constant-time operations
- the height of the segment tree is  $\log(n)$ , as it is a binary tree

**Space Complexity:**  $O(n)$

- the segment tree is represented using a binary tree structure with additional information for each node, the space required is proportional to the number of nodes in the tree

# Update

The second kind of query is of the form `update(l, r, t)` that replaces every value `a[i]` with `min(a[i], t)`, where `l <= i <= r`.

The function

```
fn update_rec(curr_node: &mut Node, l: usize, r: usize, t: u32)
```

has as parameters:

- `curr_node` , the current root of the tree
- `l` , the left endpoint of the queried segment
- `r` , the right endpoint of the queried segment

and it recursively updates the values in `a[l..r]` .

The function behaves as follows:

- if the current node segment is outside the queried update simply return
- if the queried segment is fully inside the current node's segment, update the `curr_node.key` and propagate the update on the left and right subtree recurring on `left_child` and `right_child` and then return
- recur on both the left and right subtrees of the current node and update the current node's key, which might have changed due to the update

## Problem 2: Is There

Given `n` segments `[l,r]` such that  $0 \leq l \leq r \leq n - 1$  the task is to solve `m` queries `is_there(i,j,k)` , where `is_there(i,j,k)` has to return 1 if there exists a position  $p$  with  $0 \leq i \leq p \leq j \leq n - 1$ .

The solution exploits two techniques:

- prefix sum
- binary search

The function

```
fn is_there(segments: &Vec<(u32, u32)>, query: (u32, u32, u32)) -> bool
```

has as parameters:

- `segments` , the array of segments `[l,r]` received in input
- `query` , a single triple, where the first element is `i` , the second is `j` and the last is `k`

and it returns `true` if such a position is found.

The function behaves as follows:

- create an array `axis` of length `segments.len()+1` , initialized to all zeroes

- mark the array `axis` so that each element of it represents the number of "active" segments that start in `i`
  - for each `segment` in `segments`
    - `axis[segment.0]+=1`
    - `axis[segment.1+1]-=1`
- compute the number of active segments for every position, this is done by computing the prefix sum of `axis`
- create a map that associates the number of active segments on a position `i` to the "list" of positions that have that number of active segments
- exploit binary search:
  - to the binary search pass the "list" of positions that have `k` active segments, retrieved through the map
  - search the "list" of positions, returning `true` as soon as a position that is in `[i,j]` is found.
  - if such a position is not found return `false`

**Time and Space Complexity** are analogous to the previous problem.